
Task Document: The "Agentic Architect" Sprint

Role: AI Engineer (Agentic Systems)

Time Limit: 5 Days from receipt.

Tech Stack: Python (LangGraph) + React (TypeScript) + MCP (Model Context Protocol).

1. The Mission

At Cerina, we don't just build chatbots; we build autonomous systems that act as clinical foundries. Your task is to architect and build the "**Cerina Protocol Foundry**"—an intelligent multi-agent system that autonomously designs, critiques, and refines CBT (Cognitive Behavioral Therapy) exercises.

The Twist: We are not giving you a blueprint. You must act as the **System Architect**. You decide the agent topology. You decide how they collaborate.

Constraint: This scope is deliberately massive for 3 days. You are **expected** to use AI coding assistants (Cursor, Claude, Copilot) to scaffold, generate, and debug code rapidly. We are evaluating your ability to orchestrate complex systems, not your typing speed.

2. The Requirements

A. The Backend (The "Brain")

Framework: Python & LangGraph.

Database: You must use a persistent backend (SQLite or Postgres) with LangGraph Checkpointers.

1. The Agent Architecture (Your Choice):

We need a system that mimics a rigorous clinical review board. A simple linear chain ($A \rightarrow B \rightarrow C$) is insufficient. We want to see autonomy and complex reasoning.

- *The Goal:* Produce a safe, empathetic, and structured CBT exercise based on a user intent (e.g., "*Create an exposure hierarchy for agoraphobia*").
- *The Team:* You decide the roster. However, a robust solution likely needs agents acting as **Draftsmen**, **Safety Guardians** (checking for self-harm/medical advice), **Clinical Critics** (judging tone/empathy), and perhaps a **Supervisor/Manager** to route tasks and decide when a draft is "good enough."
- *The Pattern:* Choose an architecture that best solves this (e.g., **Supervisor-Worker**, **Hierarchical Teams**, or **Network/Swarm**).
- *Autonomy:* The system should be able to loop, self-correct, and debate internally *before* disturbing the human.

2. Deep State Management ("The Blackboard"):

The agents must share a rich, structured state. It shouldn't just be a list of messages. Think of it as a shared project workspace.

- **Context:** Detailed scratchpads where agents can leave notes for each other (e.g., "Safety Agent flagged line 3; Drafter needs to revise").
- **Versions:** Ability to track previous drafts vs. current drafts.
- **Metadata:** Iteration counts, safety scores, empathy metrics.

3. Persistence & Memory:

- **Checkpointing:** Every step of the graph must be check-pointed to the database. If the server crashes, it should resume exactly where it left off.
 - **History:** The system must retain a log of all past queries and generated protocols in the database.
-

B. The Interfaces (The "Body")

1. Interface A: The React Dashboard (Human-in-the-Loop)

- **Visualization:** Build a UI that makes the "Black Box" transparent. We want to see the agents working in real-time (streaming thoughts/actions).
- **The "Halt" Mechanism:** The graph must **interrupt** execution before finalizing.
 - The UI must fetch the current state from the checkpoint.
 - It must present the generated draft to the Human User.
 - The Human can **Edit** the text or **Approve** it.
 - Only upon approval does the graph resume and save the final artifact.

2. Interface B: The MCP Server (Machine-to-Machine)

- Implement the **Model Context Protocol (MCP)** using the mcp-python SDK.
 - Expose your complex LangGraph workflow as a single **Tool** (resource) to the MCP ecosystem.
 - **Use Case:** A user on an MCP Client (like Claude Desktop) should be able to prompt: "Ask Cerina Foundry to create a sleep hygiene protocol." This triggers your backend, runs the agents, and returns the result—bypassing the React UI but using the same underlying logic.
-

3. Recommended Resources

- **Architecture:** *Building Effective Agents* (Anthropic), *LangGraph Multi-Agent Supervisor* tutorials.
 - **Protocol:** modelcontextprotocol.io (MCP Documentation).
-

4. Submission Guidelines

Deadline: 5 days.

Deliverables:

1. **Code Repo:** Modular, clean, and well-documented.
 2. **Architecture Diagram:** A visual representation of your chosen agent topology.
 3. **Loom Video (Critical - Max 5 mins):**
 - **The React UI:** Show the agents debating/refining a draft, the "Human-in-the-Loop" interruption, and the final approval.
 - **The MCP Demo:** Connect your server to a client (e.g., Claude Desktop) and trigger the workflow remotely.
 - **The Code:** Briefly explain your state definition and checkpointer logic.
-

5. Evaluation Criteria

1. **Architectural Ambition:** Did you build a trivial chain, or did you design a robust, self-correcting system?
2. **State Hygiene:** How effectively did you use the shared state/scratchpad?
3. **Persistence:** Does the Human-in-the-Loop flow work reliably using database checkpoints?
4. **MCP Integration:** Did you successfully implement the new interoperability standard?
5. **AI Leverage:** Did you use AI coding tools to deliver a "weeks worth of work" in 5 days?

Good luck. Show us what you can build.